



St. Louis Clojure

Powderkeg

Christopher Mark Gore

cgore.com

Tuesday, March 21, AD 2017

**We write Clojure at The Climate Corporation,
and we're hiring! Come work with us!**



Especially now that Bayer is buying us!



It's a pretty cool place to work, we've even got a giant globe to play with.



Clojure is a lisp.



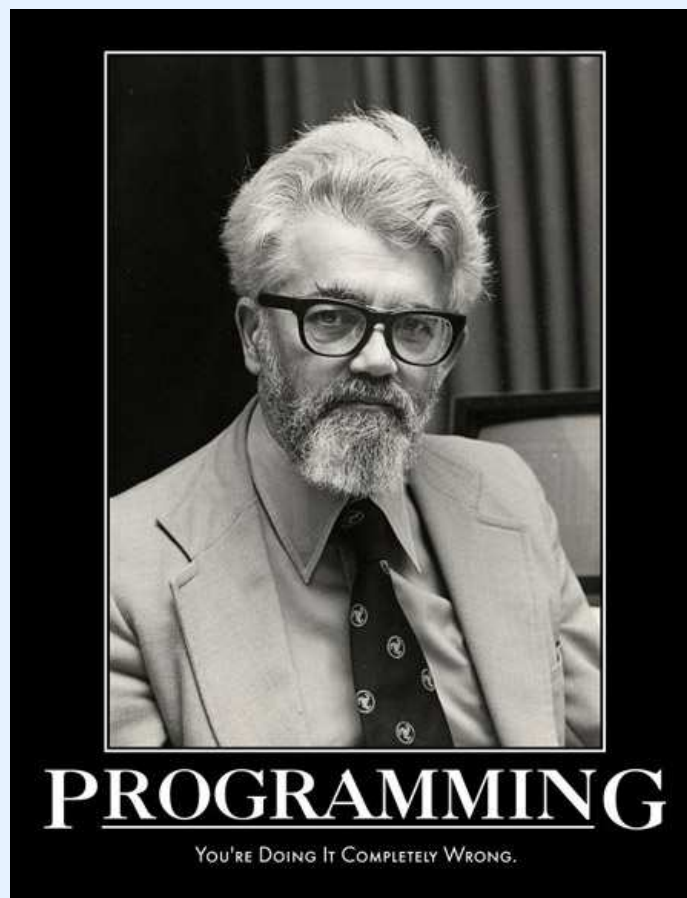
Clojure is a lisp on the JVM.



Scala is also on the JVM.



But Scala isn't a lisp.



Apache Spark is really cool, but it's in Scala.



Apache Spark is an open source cluster computing framework.

- Based on the RDD, resilient distributed dataset.
- An RDD is basically a distributed read-only multiset.
- RDDs allow for the power of MapReduce but with a lot more flexibility.
- RDDs can be treated as shared memory.
- This allows for iterative algorithms, not just map and then reduce operations.

What's a Clojurian to do?

- Use Scala? Nope.
- Here at Climate, we made `clj-spark` well before I came to work here, which was a good start.
- This eventually became Flambo, which is pretty good, but doesn't exactly feel like Clojure.

Let's make another library!

- Igor Ges and Christophe Grand introduced a new library called *Powderkeg* to work with Apache Spark in Clojure.
- It looks like normal Clojure code, thanks to transducers!
- But ...it's still really early alpha.

It looks almost like normal Clojure.

```
1 ;; 'normal' Clojure
2 (into [] (map #(* % %))
3         [1 2 3 4 5]))
4
5 ;; Flambo
6 (-> (f/parallelize sc [1 2 3 4 5])
7      (f/map (f/fn [x] (* x x)))
8      f/collect)
9
10 ;; Powderkeg
11 (into [] (keg/rdd [1 2 3 4 5]
12                  (map #(* % %)))))
```

So what exactly do you do with a Spark cluster?

- Make a big list of stuff, and RDD.
- Map on that RDD ...
- Filter down that RDD ...
- Reduce on that RDD ...
- ...

Until you have the final result of your computation. But, all this mapping, reducing, and filtering has occurred on multiple machines, not just one machine.


Why we care at Climate *(besides it just being cool.)*

- Most of our imagery generation is quite amenable to this sort of parallelization.
- We currently have to run on somewhat large (and expensive) instances because we currently operate on a per-field basis.
- The instance type required for a 100-acre field is a lot cheaper than what we need to calculate on a 3,000-acre field.
- But if we can span across multiple instances per-field, then we can use smaller and cheaper instances, just more of them.

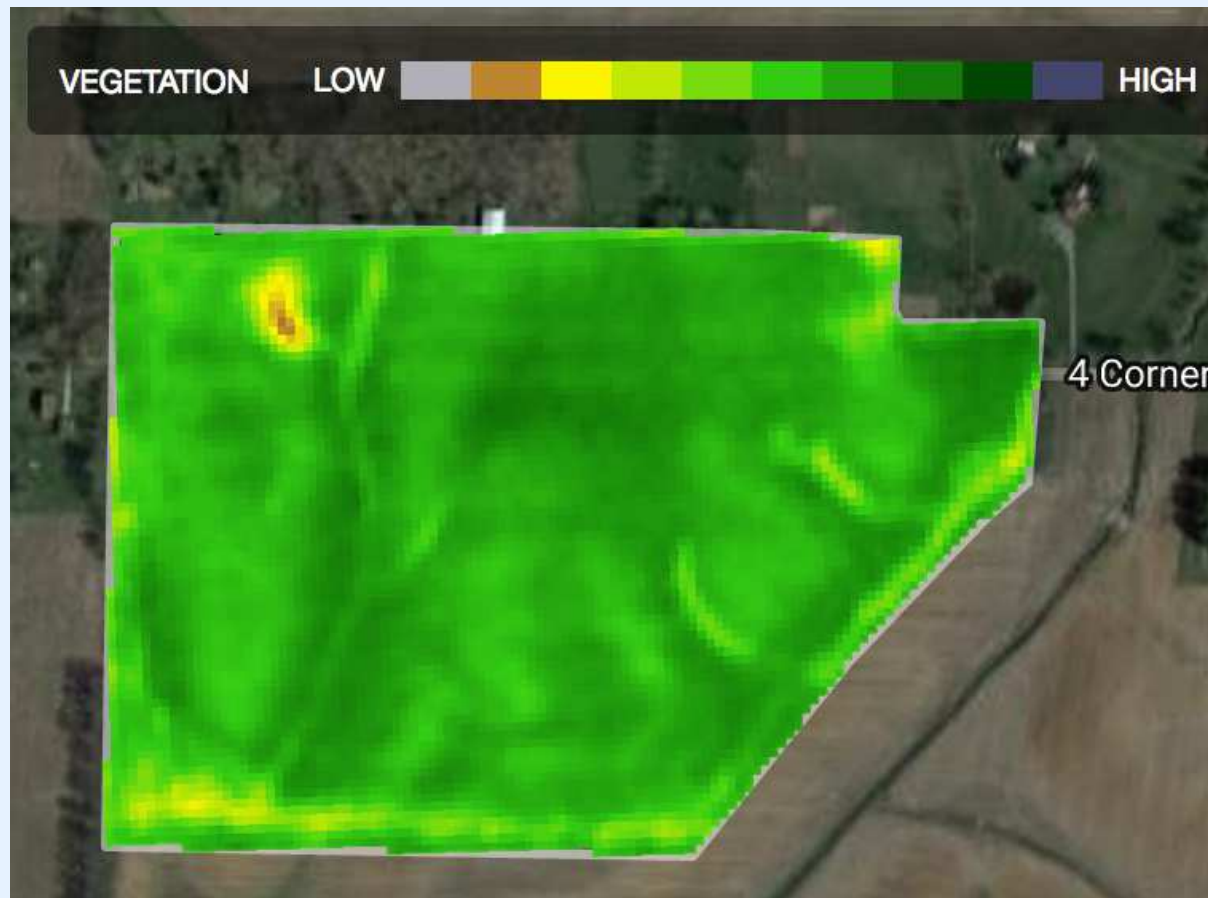
Why you care at *YourMegaCorp, Inc.*
(*besides it just being cool.*)



Actually, there's various limits to the sizes of things, so *biggish* data.



Pictures of corn ...*from space*.



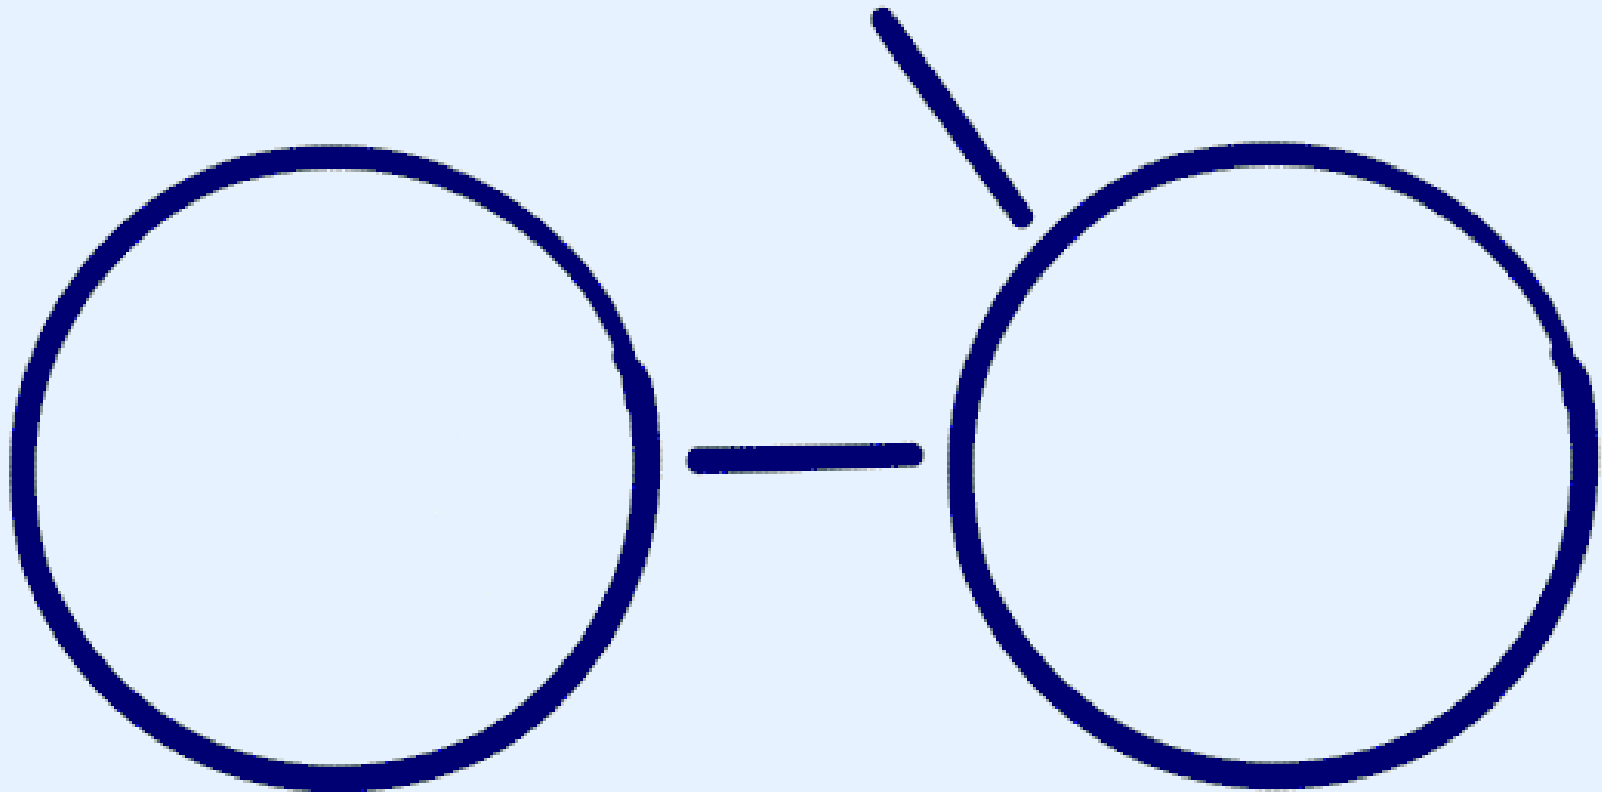
Spark SQL – there's work in progress to move this over.

<https://github.com/HCADatalab/powderkeg/tree/sql>

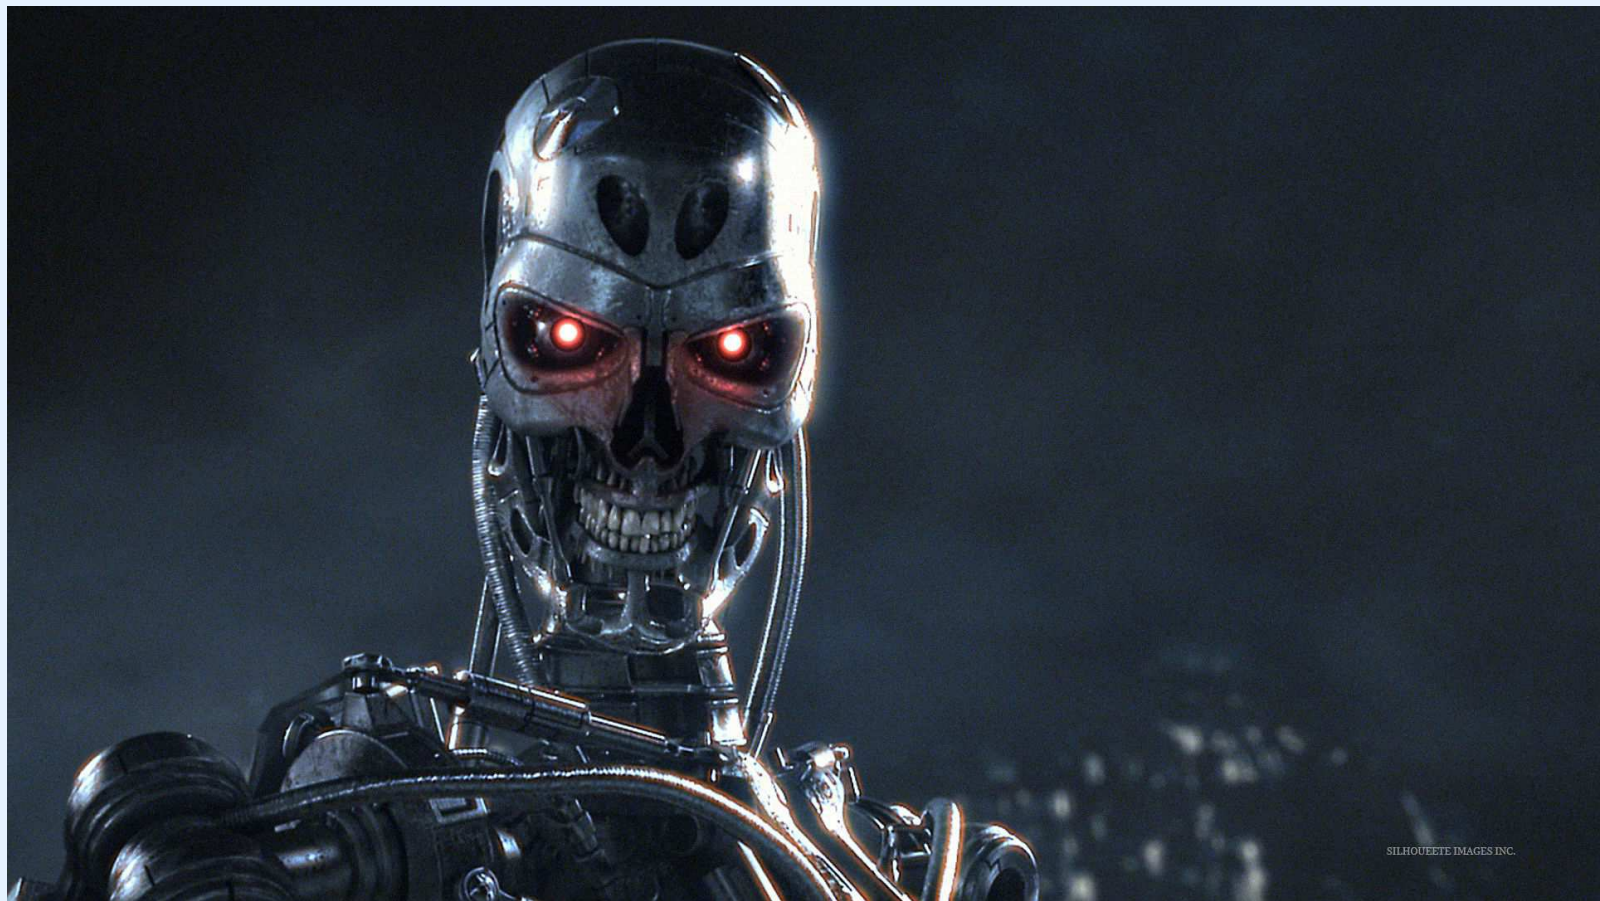
```
1 // Select customer name column
2 dfCustomers.select("name").show()
3 // Select customer name and city columns
4 dfCustomers.select("name", "city").show()
5 // Select a customer by id
6 dfCustomers.filter(
7     dfCustomers("customer_id").equalTo(500)
8 ).show()
9 // Count the customers by zip code
10 dfCustomers.groupBy("zip_code").count().show()
```

**Spark Streaming – eventually, I don't think
there's any work on that yet.**

You can implement a batch-oriented lambda architecture
in this, which is something that I need for Thinking Bicycle.



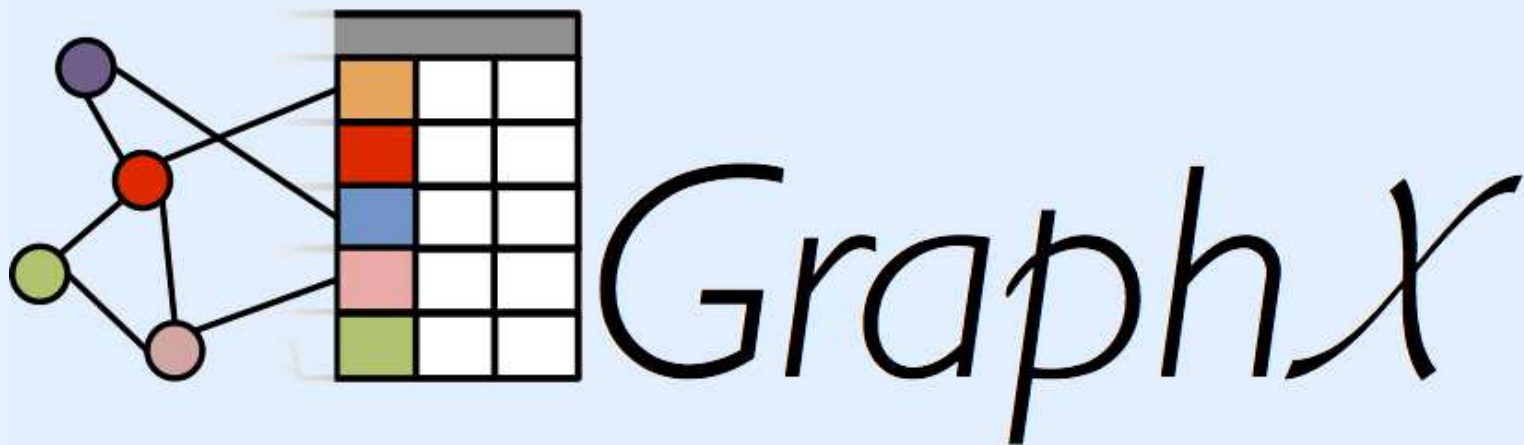
MLlib: Machine learning is cool, and there's a library for that on top of Spark. I really want it, so I'll probably make it.



MLlib's SVMs (*support vector machines*) more specifically. Because clouds suck when you've only got civilian satellites.



GraphX, a graph processing framework on top of Apache Spark, also looks pretty interesting and useful.



RDDs have *actions* which return values (1/2).

reduce aggregate the elements with some function (*think summation and the + function*)

collect return a normal array

count how many items in the RDD

first the first item in the RDD

take return a normal array of n items

takeSample random sampling of n elements

takeOrdered first n elements based on an ordering

RDDs have *actions* which return values (2/2).

saveAsTextFile write out a simple text file

saveAsSequenceFile write out a Hadoop SequenceFile

saveAsObjectFile write out a simple serialized file

countByKey for (k, v) RDDs, return (k, n) of counts

foreach run a function over the entire RDD for side effects

RDDs have *transformations* which return new RDDs (1/3).

map new RDD by mapping f

filter new RDD filtered for f is true

flatMap non-bijective map (non-1:1)

mapPartitions like map, but running separately on each partition (block) of the RDD

mapPartitionsWithIndex like *mapPartitions*, but also with an integer index of the partition

sample random sample of the RDD

RDDs have *transformations* which return new RDDs (2/3).

union union of two RDDs

intersection intersection of two RDDs

distinct the distinct (unique) elements as a new RDD

groupByKey (k, v) to $(k, \{v_1, v_2, \dots\})$

reduceByKey like reduce, but on a per-key basis

aggregateByKey per-key merge of the RDD

sortByKey an RDD sorted by the keys

RDDs have *transformations* which return new RDDs (3/3).

join (k, v) and (k, w) become $(k, (v, w))$

cogroup groups two RDDs together on their shared keys

cartesian this is the reason why some people shouldn't be allowed near databases

pipe pipe through a shell command

coalesce reduce the number of partitions (blocks)

repartition set the number of partitions (blocks)

repartitionAndSortWithinPartitions more efficient than a repartition and then sort

Questions?